Processing Big Data with Hadoop in Azure HDInsight

Lab 2 – Processing Big Data with Hive

Overview

In this lab, you will process data in web server log files by creating Hive tables, populating them with data, and querying them.

What You'll Need

To complete the labs, you will need the following:

- A web browser
- A Microsoft account
- A Microsoft Azure subscription
- A Microsoft Windows, Linux, or Apple Mac OS X computer on which the Azure CLI has been installed.
- The lab files for this course

Note: To set up the required environment for the lab, follow the instructions in the **Setup** document for this course.

Preparing for Data Processing

Before you can use Hive to analyze the log files, you must provision an HDInsight cluster and upload the source data to Azure storage.

Provision an HDInsight Cluster

Note: If you already have an HDInsight cluster and associated storage account, you can skip this task.

1. Use the steps provided in Lab 1 to provision a Hadoop HDInsight cluster on Linux and an associated storage account.
2. Make a note of the details of your HDInsight cluster configuration.

View Source Data

1. In the **HDILabs** folder where you extracted the lab files, in the **Lab02\iislogs** folder, open any of the text files in a text editor. Each file in this folder is an Internet Information Services (IIS) log file

from a web server, and there is a file for each month between January and June 2008. When you have viewed the information in the file, close it without saving any changes.

2. View the contents of the **Lab02\iislogs_gz** folder. This contains the same log files compressed using the gzip compression algorithm, significantly reducing the size of the files.

Upload the Log Files to Azure Storage

Now you are ready to upload the source data to Azure storage for processing. The instructions here assume you will use the Azure CLI to do this, but you can use any Azure Storage tool you prefer.

1. Open a new command line console on your local computer.
2. Enter the following command to switch the Azure CLI to *resource manager* mode.

```
azure config mode arm
```

> **Note**: If a *command not found* error is displayed, ensure that you have followed the instructions in the setup guide to install the Azure CLI.

3. Enter the following command to log into your Azure subscription:

```
azure login
```

4. Follow the instructions that are displayed to browse to the Azure device login site and enter the authentication code provided. Then sign into your Azure subscription using your Microsoft account.
5. Enter the following command to view your Azure resources:

```
azure resource list
```

6. Verify that your HDInsight cluster and the related storage account are both listed. Note that the information provided includes the resource group name as well as the individual resource names.
7. Enter the following command on a single line to determine the connection string for your Azure storage account, specifying the storage account and resource group names you noted earlier:

```
azure storage account connectionstring show storage_account -g
resource_group
```

8. Note the connection string, copying it to the clipboard if your command line tool supports it.
9. If you are working on a Windows client computer, enter the following command to set a system variable for the connection string:

```
SET AZURE_STORAGE_CONNECTION_STRING=your_connection_string
```

If you are using a Linux or Mac OS X client computer, enter the following command to set a system variable for the connection string (enter the connection string in quotation marks):

```
export AZURE_STORAGE_CONNECTION_STRING="your_connection_string"
```

10. Enter the following command on a single line to upload the **2008-01.txt.gz** file to a blob named **data/logs/2008-01.txt.gz** in the container used by your HDInsight cluster. Replace *local_path* with the local path to 2008-01.txt.gz (for example *c:\HDILabs\Lab02\iislogs_gz\2008-01.txt.gz* or *HDILabs/Lab02/iislogs_gz/2008-01.txt.gz*) and replace *container* with the name of the default storage container used by your cluster:

```
azure storage blob upload local_path container data/logs/2008-01.txt.gz
```

11. Wait for the file to be uploaded.
12. Repeat the previous step to upload each of the remaining **2008-0X.txt.gz** log files in the local **iislogs_gz** folder to blobs named **data/logs/2008-0X.txt.gz** (where *X* is the month number). There should be six files in total.

Creating, Loading, and Querying Hive Tables

Now that you have provisioned an HDInsight cluster and uploaded the source data, you can create Hive tables and use them to process the data. To do this, you will need to open an SSH console that is connected to your cluster.

If you are using a Windows client computer:

1. In the Microsoft Azure portal, on the **HDInsight Cluster** blade for your HDInsight cluster, click **Secure Shell**, and then in the **Secure Shell** blade, under **Windows users**, copy the **Host name** (which should be *your_cluster_name*-**ssh.azurehdinsight.net**) to the clipboard.
2. Open PuTTY, and in the **Session** page, paste the host name into the **Host Name** box. Then under **Connection type**, select **SSH** and click **Open**. If a security warning that the host certificate cannot be verified is displayed, click **Yes** to continue.
3. When prompted, enter the SSH username and password you specified when provisioning the cluster (not the cluster login).

If you are using a Mac OS X or Linux client computer:

1. In the Microsoft Azure portal, on the **HDInsight Cluster** blade for your HDInsight cluster, click **Secure Shell**, and then in the **Secure Shell** blade, under **Linux, Unix, and OS X users**, note the command used to connect to the head node.
2. Open a new terminal session, and enter the appropriate command to connect, specifying your SSH user name (not the cluster login) and cluster name as necessary:

   ```
   ssh your_ssh_user_name@your_cluster_name-ssh.azurehdinsight.net
   ```

3. If you are prompted to connect even though the certificate can't be verified, enter **yes**.
4. When prompted, enter the password for the SSH username.

Note: If you have previously connected to a cluster with the same name, the certificate for the older cluster will still be stored and a connection may be denied because the new certificate does not match the stored certificate. You can delete the old certificate by using the **ssh-keygen** command, specifying the path of your certificate file (**f**) and the host record to be removed (**R**) - for example:

```
ssh-keygen -f "/home/usr/.ssh/known_hosts" -R clstr-ssh.azurehdinsight.net
```

View the Uploaded Data Files

Now that you have a remote command console for your cluster, you can verify that the log files have been uploaded to the shared cluster storage.

1. In the SSH console for your cluster, enter the following command to view the contents of the **/data/logs** folder in the HDFS file system.

   ```
   hdfs dfs -ls /data/logs
   ```

2. Verify that the folder contains six log files.

Create a Hive Table for the Raw Log Data

While there are many Hive editors available, and you can use any one that you prefer. The instructions here assume that you will use the Hive command line interface to work with Hive in this lab.

1. In the SSH console for your cluster, enter the following command to start the Hive command line interface:

   ```
   hive
   ```

2. In the Hive command line interface, enter the following HiveQL statement to create a table named **rawlogs** (you can copy and paste this from **Create Raw Table.txt** in the **HDILabs\Lab02** folder):

   ```
   CREATE TABLE rawlog
   (log_date STRING,
    log_time STRING,
    c_ip STRING,
    cs_username STRING,
    s_ip STRING,
    s_port STRING,
    cs_method STRING,
    cs_uri_stem STRING,
    cs_uri_query STRING,
    sc_status STRING,
    sc_bytes INT,
    cs_bytes INT,
    time_taken INT,
    cs_user_agent STRING,
    cs_referrer STRING)
   ROW FORMAT DELIMITED FIELDS TERMINATED BY ' ';
   ```

3. Wait for the job to complete.
4. Enter the following command to verify that the table has been created (there may also be an existing table named **hivesampletable**):

   ```
   SHOW TABLES;
   ```

5. Enter the following command to query the table, and verify that no rows are returned:

   ```
   SELECT * FROM rawlog;
   ```

Load the Source Data into the Raw Log Table

1. In the Hive command line interface, enter the following HiveQL statement to move the log files you previously uploaded into the folder for the **rawlogs** table (you can copy and paste this from **Load Raw Data.txt** in the **HDILabs\Lab02** folder):

   ```
   LOAD DATA INPATH '/data/logs' INTO TABLE rawlog;
   ```

2. Wait for the job to complete.
3. Enter the following command to return the first 100 rows in the **rawlogs** table:

   ```
   SELECT * FROM rawlog LIMIT 100;
   ```

4. View the output, noting that the query retrieved the first 100 rows from the table in tab-delimited text format. Note that the output includes rows containing comments from the source document (the first column value for these rows is prefixed with a **#** character, and null values are used for columns in the table schema for which there are no values in the source data).

Clean the Log Data

1. In the Hive command line interface, enter the following HiveQL statement to create an external table named **cleanlog** based on the **/data/cleanlog** folder (you can copy and paste this from **Create Clean Table.txt** in the **HDILabs\Lab02** folder):

```
CREATE EXTERNAL TABLE cleanlog
(log_date DATE,
 log_time STRING,
 c_ip STRING,
 cs_username STRING,
 s_ip STRING,
 s_port STRING,
 cs_method STRING,
 cs_uri_stem STRING,
 cs_uri_query STRING,
 sc_status STRING,
 sc_bytes INT,
 cs_bytes INT,
 time_taken INT,
 cs_user_agent STRING,
 cs_referrer STRING)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ' '
STORED AS TEXTFILE LOCATION '/data/cleanlog';
```

2. Wait for the job to complete.
3. Enter the following command to extract rows that do not contain comments from the **rawlogs** table, and insert them into the **cleanlog** table (you can copy and paste this from **Load Clean Data.txt** in the **HDILabs\Lab02** folder):

```
INSERT INTO TABLE cleanlog
SELECT *
FROM rawlog
WHERE SUBSTR(log_date, 1, 1) <> '#';
```

4. Wait for the job to complete.
5. Enter the following command to retrieve the first 100 rows from the **cleanlog** table:

```
SELECT * FROM cleanlog LIMIT 100;
```

6. View the results returned by the query, noting that the **cleanlog** table does not include any rows prefixed with a **#** character.

Create a View

1. In the Hive command line interface, enter the following command to create a view named **vDailySummary** (you can copy and paste this from **Create View.txt** in the **HDILabs\Lab02** folder):

```
CREATE VIEW vDailySummary
AS
SELECT log_date,
       COUNT(*) AS requests,
       SUM(cs_bytes) AS inbound_bytes,
       SUM(sc_bytes) AS outbound_bytes
FROM cleanlog
GROUP BY log_date;
```

2. Wait for the job to complete.
3. Enter the following command to query the **vDailySummary** view:

```
SELECT * FROM vDailySummary
ORDER BY log_date;
```

4. Wait for the job to complete and then view the output returned by the query.

Partitioning Data

Partitioning data can improve performance when queries commonly filter on specific columns, or when the job can be effectively split across multiple reducer nodes.

Create and Load a Partitioned Table

1. In the Hive command line interface, enter the following command to create a partitioned table (you can copy and paste this from **Create Partitioned Table.txt** in the **HDILabs\Lab02** folder):

```
CREATE EXTERNAL TABLE partitionedlog
(log_day int,
 log_time STRING,
 c_ip STRING,
 cs_username STRING,
 s_ip STRING,
 s_port STRING,
 cs_method STRING,
 cs_uri_stem STRING,
 cs_uri_query STRING,
 sc_status STRING,
 sc_bytes INT,
 cs_bytes INT,
 time_taken INT,
 cs_user_agent STRING,
 cs_referrer STRING)
PARTITIONED BY (log_year int, log_month int)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ' '
STORED AS TEXTFILE LOCATION '/data/partitionedlog';
```

2. Enter the following command to load data from the **rawlog** table into the partitioned table (you can copy and paste this from **Load Partitioned Table.txt** in the **HDILabs\Lab02** folder):

```
SET hive.exec.dynamic.partition.mode=nonstrict;
SET hive.exec.dynamic.partition = true;
INSERT INTO TABLE partitionedlog PARTITION(log_year, log_month)
SELECT DAY(log_date),
       log_time,
       c_ip,
       cs_username,
       s_ip,
       s_port,
       cs_method,
       cs_uri_stem,
       cs_uri_query,
       sc_status,
       sc_bytes,
       cs_bytes,
       time_taken,
       cs_user_agent,
       cs_referrer,
       YEAR(log_date),
       MONTH(log_date)
FROM rawlog
WHERE SUBSTR(log_date, 1, 1) <> '#';
```

3. Wait for the query job complete. By partitioning the data in this way, queries filtered on year and month will benefit from having to search a smaller volume of data to retrieve the results.
4. Enter the following command to query the partitioned table, and when the job has completed, view the results:

```
SELECT log_day, count(*) AS page_hits
FROM partitionedlog
WHERE log_year=2008 AND log_month=6
GROUP BY log_day;
```

View Folders for a Partitioned Table

1. In the Hive command line interface, enter the following command to exit Hive:

```
exit;
```

2. At the command prompt, enter the following command to view the **/data/partitionedlog** folder:

```
hdfs dfs -ls /data/partitionedlog
```

3. Note that the folder contains a subfolder for each value in the first partitioning key (in this case, a single folder named **log_year=2008**).
4. Enter the following command to view the contents of the **log_year=2008** subfolder (note that quotation marks are required because the folder name includes a "=" character):

```
hdfs dfs -ls "/data/partitionedlog/log_year=2008"
```

5. Note that the subfolder contains a further subfolder for each value in the second partitioning key (in this case, a folder for each **log_month** value in the source data).
6. Enter the following command to view the contents of the **log_month=6** subfolder:

```
hdfs dfs -ls "/data/partitionedlog/log_year=2008/log_month=6"
```

7. Note that the subfolders for each month contain one or more data files. In this case, a file named **000000_0** contains all of the log records with a year of 2008 and a month of 6.
8. Enter the following command to view the last few kilobytes of data in the file:

```
hdfs dfs -tail "/data/partitionedlog/log_year=2008/log_month=6/000000_0"
```

9. Note that the first field in the data file contains the day of the month (the last few rows will contain the value **30**.) The year and month are not stored in the data file, but are determined from the partitioning subfolder in which the file is stored - this file contains only the records for June 2008.
10. Close all open command windows.

Cleaning Up

Now that you have finished this lab, you can delete the HDInsight cluster and storage account. This ensures that you avoid being charged for cluster resources when you are not using them. If you are using a trial Azure subscription that includes a limited free credit value, deleting the cluster maximizes your credit and helps to prevent using it all before the free trial period has ended.

Note: If you are proceeding straight to the next lab, omit this task and use the same cluster in the next lab. Otherwise, follow the steps below to delete your cluster and storage account.

Delete Cluster Resources

1. In the Microsoft Azure portal, click **Resource Groups**.

2. On the **Resource groups** blade, click the resource group that contains your HDInsight cluster, and then on the **Resource group** blade, click **Delete**. On the confirmation blade, type the name of your resource group, and click **Delete**.
3. Wait for your resource group to be deleted, and then click **All Resources**, and verify that the cluster, and the storage account that was created with your cluster, have both been removed.
4. Close the browser.